

CS6375: Machine Learning

Gautam Kunapuli

Decision Trees



THE UNIVERSITY OF TEXAS AT DALLAS

Erik Jonsson School of Engineering and Computer Science

Example: Restaurant Recommendation

Example: Develop a model to **recommend restaurants** to users depending on their past dining experiences.

Here, the features are **cost** (x_1) and the user's **spiciness rating** of the food at the restaurant (x_2) and the label is if they liked the food ($y_i = \text{green circle}$) or not ($y_i = \text{red square}$).

A data set is **linearly separable** if there exists a hyperplane that separates positive examples from negative examples.


- Relatively easy to learn (using standard techniques)
- Easy to visualize and interpret

Many **data sets in real world** are **not linearly separable**!

Two options:

- Use **non-linear features**, and learn a linear classifier in the transformed non-linear feature space
- Use **non-linear classifiers**

Decision Trees can handle nonlinear separable data sets and are one of the **most popular classifiers**



Decision Trees: Introduction

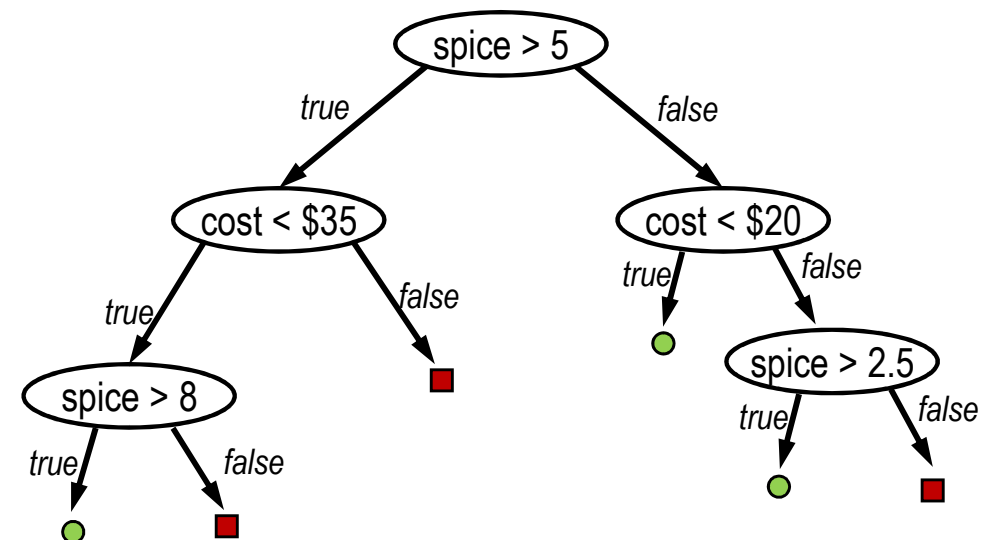

Example: Develop a model to **recommend restaurants** to users depending on their past dining experiences.

Here, the features are **cost** (x_1) and the user's **spiciness rating** of the food at the restaurant (x_2) and the label is if they liked the food ($y_i = \text{green circle}$) or not ($y_i = \text{red square}$).

Decision Trees represent decision-making as a **checklist of questions**, and visualize it using a tree-structure

Decision Tree **representation**:

- Each **non-leaf node** tests an **attribute/feature**
- Each **branch** corresponds to **attribute/feature** value, a decision (to choose a path) as a result of the test
- Each **leaf node** assigns a **classification**

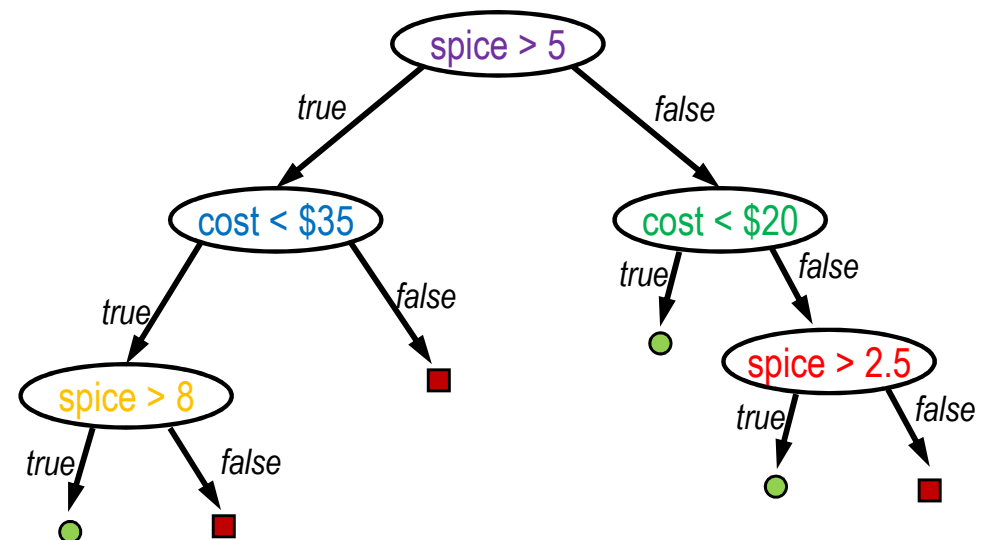



Decision Trees: Introduction

Example: Develop a model to **recommend restaurants** to users depending on their past dining experiences.

Here, the features are **cost** (x_1) and the user's **spiciness rating** of the food at the restaurant (x_2) and the label is if they liked the food ($y_i = \text{green circle}$) or not ($y_i = \text{red square}$).

- Decision trees divide the feature space into **axis-parallel rectangles**
- Decision Trees can handle **arbitrarily non-linear representations**, given sufficient tree complexity
- Worst-case scenario: the decision tree has an **exponential number of nodes!** (why?)



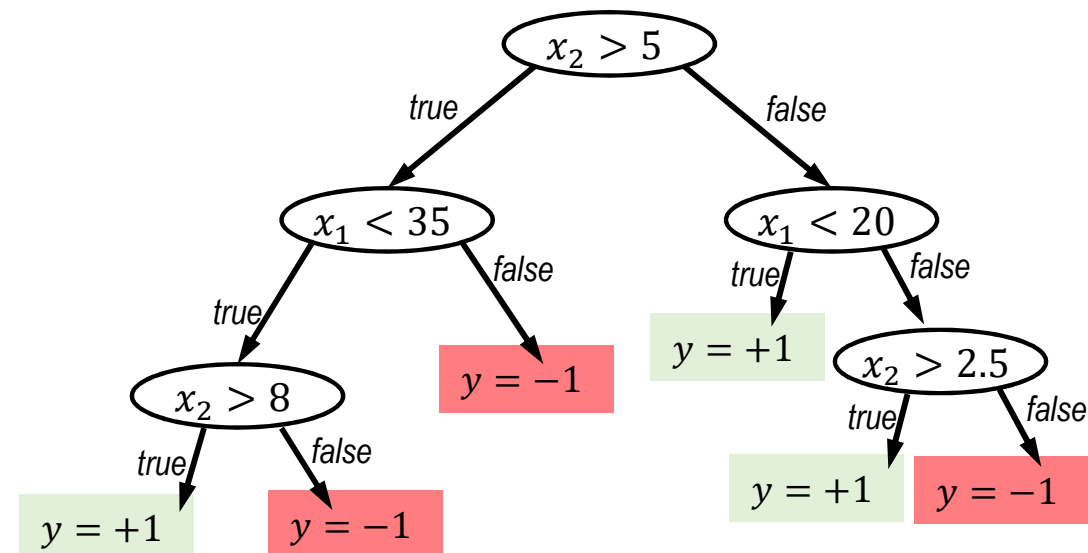

Decision Trees: Introduction

Example: Develop a model to **recommend restaurants** to users depending on their past dining experiences.

Here, the features are **cost** (x_1) and the user's **spiciness rating** of the food at the restaurant (x_2) and the label is if they liked the food ($y = +1$) or not ($y = -1$).

- Decision trees divide the feature space into **axis-parallel rectangles**
- Decision Trees can handle **arbitrarily non-linear representations**, given sufficient tree complexity
- Worst-case scenario: the decision tree has an **exponential number of nodes!**
 - If the target function has n Boolean features, there are 2^n possible inputs
 - In the worst case, there is one leaf node for each input (for example: XOR)


Decision trees are not unique, and many decision trees can represent the same hypothesis!



Decision Trees: Introduction

Example: Develop a model to **recommend restaurants** to users depending on their past dining experiences.

Here, the features are **cost** (x_1) and the user's **spiciness rating** of the food at the restaurant (x_2) and the label is if they liked the food ($y = +1$) or not ($y = -1$).




When do you want Decision Trees?

When instances are **describable by attribute-value pairs**:

- target function is **discrete-valued**
- **disjunctive hypothesis** may be required
- need for **interpretable** model

Examples:

- Equipment or medical diagnosis
- Credit risk analysis
- Modeling calendar scheduling preferences



Learning Decision Trees

Problem Formulation: Find a decision tree h that achieves minimum misclassification errors on the training data

- **Solution Approach 1 (Naïve solution):** Create a decision tree with one path from root to leaf for each training example. *Such a tree would just memorize the training data, and will **not generalize well to new points**.*
- **Solution Approach 2 (Exact solution):** Find the **smallest** tree that minimizes the classification error. *Finding this solution is **NP-Hard**!*
- **Solution Approach 3 (Heuristic solution):** Top-down greedy search

Initialize: Choose the best feature f^* for the root of the tree

Function GrowTree(data, f^*)

¹Separate data into subsets $\{S_1, S_2, \dots, S_k\}$, where each subset S_i contains examples that have the **same value for f^***

²for $S_i \in \{S_1, S_2, \dots, S_k\}$

Choose the best feature f_i^* for the next node

Recursively GrowTree(S_i, f_i^*) until all examples have the same class label

Learning Decision Trees

Problem Formulation: Find a decision tree h that achieves minimum misclassification errors on the training data

- **Solution Approach 1 (Naïve solution):** Create a decision tree with one path from root to leaf for each training example. *Such a tree would just memorize the training data, and will **not generalize well to new points**.*
- **Solution Approach 2 (Exact solution):** Find the **smallest** tree that minimizes the classification error. *Finding this solution is **NP-Hard!***
- **Solution Approach 3 (Heuristic solution):** Top-down greedy search

Initialize: Choose the best feature f^* for the root of the tree

Function GrowTree(data, f^*)

¹Separate data into subsets $\{S_1, S_2, \dots, S_k\}$, where each subset S_i contains examples that have the **same value for f^***

²for $S_i \in \{S_1, S_2, \dots, S_k\}$

Choose **the best feature** f_i^* for the next node

Recursively GrowTree(S_i, f_i^*) until all **examples have the same class label**

How do we pick the best feature?

How do we decide when to stop?

Learning Decision Trees

Problem Formulation: Find a decision tree h that achieves minimum misclassification errors on the training data

Solution Approach 3 (Heuristic solution): Top-down greedy search

Initialize: Choose the best feature f^* for the root of the tree

Function GrowTree(data, f^*)

¹Separate data into subsets $\{S_1, S_2, \dots, S_k\}$, where each subset S_i contains examples that have the **same value for f^***

²for $S_i \in \{S_1, S_2, \dots, S_k\}$

Choose **the best feature** f_i^* for the next node




Recursively GrowTree(S_i, f_i^*) until all **examples have the same class label**

How do we pick the next best feature to place in a decision tree?

- Random choice
- Largest number of values
- Fewest number of values
- **Lowest classification error**
- Information theoretic measure (Quinlan's approach)

x_1	x_2	x_3	y
0	0	0	1
0	0	1	0
0	1	0	1
0	1	1	1
1	0	0	0
1	0	1	1
1	1	0	0
1	1	1	0

Training examples



Learning Decision Trees

Problem Formulation: Find a decision tree h that achieves minimum misclassification errors on the training data

Solution Approach 3 (Heuristic solution): Top-down greedy search

Initialize: Choose the best feature f^* for the root of the tree

Function GrowTree(data, f^*)

¹Separate data into subsets $\{S_1, S_2, \dots, S_k\}$, where each subset S_i contains examples that have the **same value for f^***

²for $S_i \in \{S_1, S_2, \dots, S_k\}$

Choose **the best feature** f_i^* for the next node




Recursively GrowTree(S_i, f_i^*) until all **examples have the same class label**

How do we pick the next best feature to place in a decision tree?



- Random choice
- Largest number of values
- Fewest number of values
- **Lowest classification error**
- Information theoretic measure (Quinlan's approach)

x_1	x_2	x_3	y
0	0	0	1
0	0	1	0
0	1	0	1
0	1	1	1
1	0	0	0
1	0	1	1
1	1	0	0
1	1	1	0

Training examples



Can think of counts as **probability distributions over the labels**



Learning Decision Trees

Problem Formulation: Find a decision tree h that achieves minimum misclassification errors on the training data


Solution Approach 3 (Heuristic solution): Top-down greedy search
Initialize: Choose the best feature f^* for the root of the tree
Function GrowTree(data, f^*)
 ¹Separate data into subsets $\{S_1, S_2, \dots, S_k\}$, where each subset S_i contains examples that have the **same value for f^***
 ²for $S_i \in \{S_1, S_2, \dots, S_k\}$
 Choose **the best feature** f_i^* for the next node
 Recursively GrowTree(S_i, f_i^*) until all **examples have the same class label**

How do we pick the next best feature to place in a decision tree?

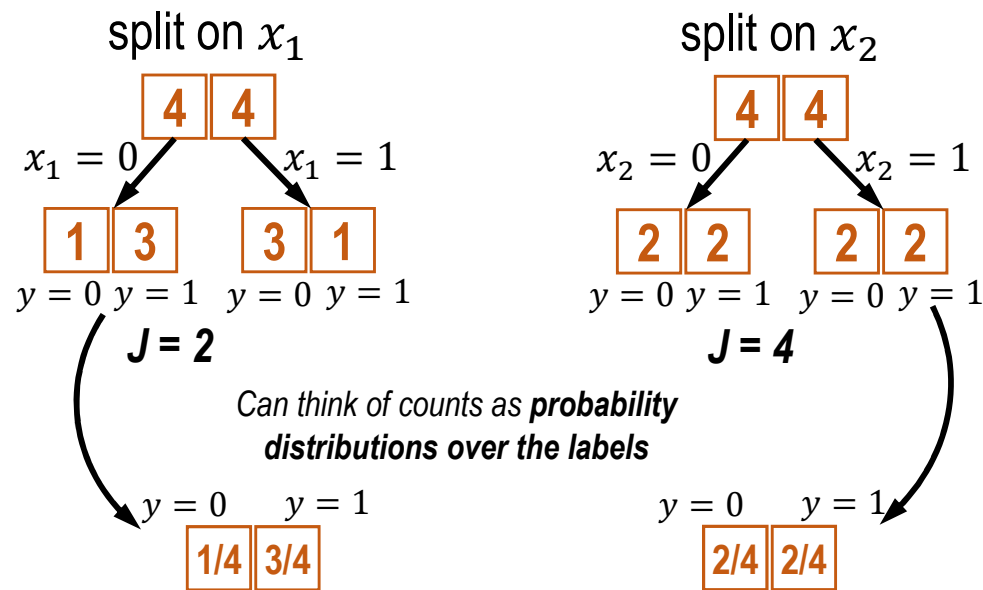

- Random choice
- Largest number of values
- Fewest number of values
- **Lowest classification error**
- Information theoretic measure (Quinlan's approach)

The selected attribute is a **good split** if we are **more “certain”** about the classification after the split (compare with the perceptron)

- If each partition with respect to the chosen attribute has a **distinct class label**, we are **completely certain** about the classification



- If **class labels are evenly divided** between partitions, we are **very uncertain** about the classification




Learning Decision Trees

Problem Formulation: Find a decision tree h that achieves minimum misclassification errors on the training data

Solution Approach 3 (Heuristic solution): Top-down greedy search
Initialize: Choose the best feature f^* for the root of the tree
Function GrowTree(data, f^*)
 ¹Separate data into subsets $\{S_1, S_2, \dots, S_k\}$, where each subset S_i contains examples that have the **same value for f^***
 ²for $S_i \in \{S_1, S_2, \dots, S_k\}$
 Choose **the best feature** f_i^* for the next node
 Recursively GrowTree(S_i, f_i^*) until all **examples have the same class label**

We need a better way to resolve the uncertainty!



How do we pick the next best feature to place in a decision tree?


- Random choice
- Largest number of values
- Fewest number of values
- **Lowest classification error**
- Information theoretic measure (Quinlan's approach)

The selected attribute is a **good split** if we are **more “certain”** about the classification after the split (compare with the perceptron)

- If each partition with respect to the chosen attribute has a **distinct class label**, we are **completely certain** about the classification



- If **class labels are evenly divided** between partitions, we are **very uncertain** about the classification



Discrete Probability and Information Theory

A **discrete probability distribution** describes the probability of occurrence of each value of a discrete random variable.

The **surprise** or **self-information** of each event of X is defined to be

$$S(X = x) = -\log_2 \text{Prob}(X = x)$$

- An event with probability 1 has zero surprise; *this is because when the content of a message is known beforehand with certainty, there is no actual information conveyed*
- The **smaller the probability** of event, the **larger the quantity of self-information** associated with the message that the event occurred
- An event with probability 0 has infinite surprise
- The surprise is the **asymptotic number of bits of information** that need to be transmitted to a recipient who knows the probabilities of the results. This is also called the **description length** of X .

Random Variable: Number of heads when tossing a coin 3 times

X	0	1	2	3
Prob(X)	1/8	3/8	3/8	1/8
$-\log_2 P(X)$	3	1.415	1.415	3
$-\log_e P(X)$	2.079	0.980	0.980	2.079
$-\log_{10} P(X)$	0.903	0.426	0.426	0.903

If the logarithm is base 2, the unit of information is bits, base e is nats and base 10 hartleys

Entropy




A standard way to measure **uncertainty of a random variable** is to use **entropy**

$$H(X) = - \sum_x P(X = x) \log_2 P(X = x)$$

- Note that the entropy is computed by **summing over all the events/outcomes/states** of the random variable.
- Entropy is maximized for uniform distributions, where the probability of all outcomes is equal (is this what we want?)
- Entropy is minimized for distributions that place all their probability on a single outcome (or is this what we want?)

The entropy of label distributions can be computed as:

$$H(y) = -P(y = 0) \log_2 P(y = 0) - P(y = 1) \log_2 P(y = 1)$$



Conditional Entropy and Mutual Information

Entropy can also be computed when conditioned on another variable:

$$H(Y|X) = - \sum_x P(X = x) \sum_y P(Y = y | X = x) \log_2 (Y = y | X = x)$$

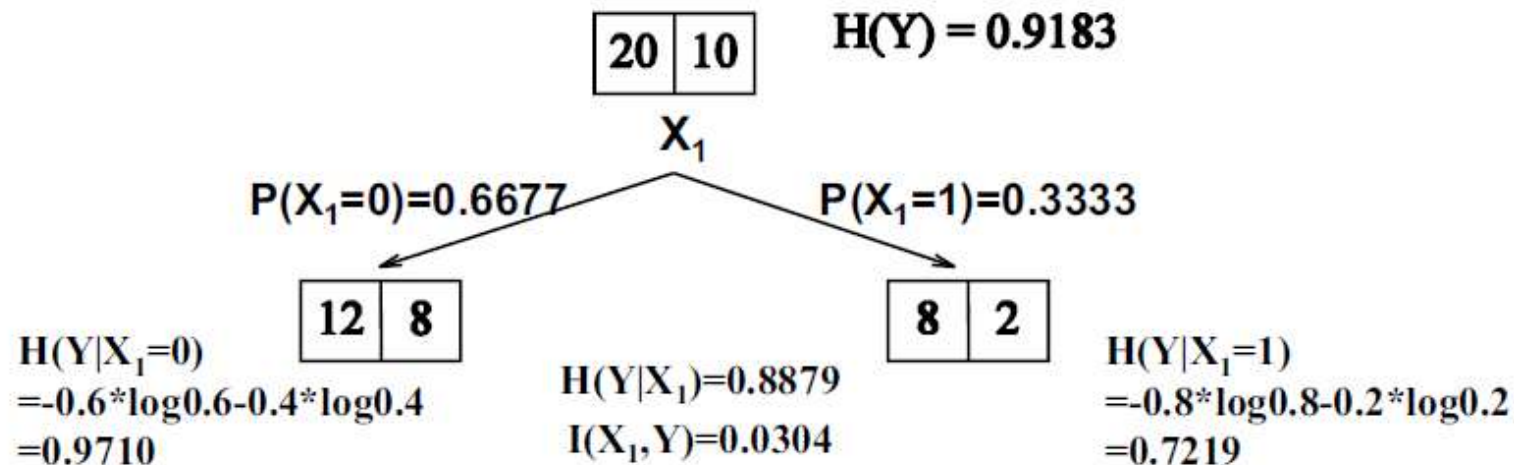
This is called **conditional entropy** and is the amount of information needed to quantify the random variable Y given the random variable X .

The **mutual information** or **information gain** between two random variables is defined as

$$I(X, Y) = H(Y) - H(Y|X)$$



This is the amount of information we learn about Y by knowing the value of X and vice-versa (it is symmetric).

In our case, **larger information gain** corresponds to **less uncertainty about Y (labels) given X (data)**.




Choosing the Best Feature

Step 1: Count the various combinations of features and labels





Step 2: Convert to probabilities



x_1	x_2	y
1	1	0 (+)
1	0	0 (+)
1	1	0 (+)
1	0	0 (+)
0	1	0 (+)
0	0	1 (-)
0	1	1 (-)
0	0	1 (-)

Choosing the Best Feature

Step 3: Compute information gain for both splits and pick the variable with the biggest gain



x_1	x_2	y
1	1	0 (+)
1	0	0 (+)
1	1	0 (+)
1	0	0 (+)
0	1	0 (+)
0	0	1 (-)
0	1	1 (-)
0	0	1 (-)

The ID3 Algorithm

The ID3 (Iterative Dichotomizer) and its successor, C4.5 were developed by Ross Quinlan in the early to mid 1980s and are widely considered to be a landmark machine learning algorithms, and until at least 2008, were the #1 data mining tool.

ID3(*Examples*, *Target_attribute*, *Attributes*)

Examples are the training examples. *Target_attribute* is the attribute whose value is to be predicted by the tree. *Attributes* is a list of other attributes that may be tested by the learned decision tree. Returns a decision tree that correctly classifies the given *Examples*.

- Create a *Root* node for the tree
- If all *Examples* are positive, Return the single-node tree *Root*, with label = +
- If all *Examples* are negative, Return the single-node tree *Root*, with label = -
- If *Attributes* is empty, Return the single-node tree *Root*, with label = most common value of *Target_attribute* in *Examples*
- Otherwise Begin
 - $A \leftarrow$ the attribute from *Attributes* that best* classifies *Examples*
 - The decision attribute for *Root* $\leftarrow A$
 - For each possible value, v_i , of A ,
 - Add a new tree branch below *Root*, corresponding to the test $A = v_i$
 - Let $Examples_{v_i}$ be the subset of *Examples* that have value v_i for A
 - If $Examples_{v_i}$ is empty
 - Then below this new branch add a leaf node with label = most common value of *Target_attribute* in *Examples*
 - Else below this new branch add the subtree
 $ID3(Examples_{v_i}, Target_attribute, Attributes - \{A\})$
- End
- Return *Root*

Some Final Details

When do we terminate?

- If the current set is “**pure**” (i.e., has a single label in the output), stop
- If you **run out of attributes to recurse on**, even if the current data set isn’t pure, stop and use a majority vote
- If a partition contains no data points, use the majority vote at its parent in the tree
- If a partition contains no data items, nothing to recurse on
- For fixed depth decision trees, the **final label is determined by majority vote**

How do we handle real-valued features?

- For continuous attributes, use threshold splits
- Split the tree into $x_k < t$ and $x_k \geq t$
- Can split on the same attribute multiple times on the same path down the tree

How do we select the splitting threshold?

Overfitting in Decision Trees

Hypothesis space is complete! *Target function is surely in there*



No back tracking; *Greedy thus local minima*

Statistics-based search choices; *Robust to noisy data*

Inductive bias: heuristically prefer shortest tree

Decision trees will always overfit!

It is always possible to obtain zero training error on the input data with a deep enough tree (if there is no noise in the labels)



Possibly just noise, but the tree is grown larger to capture these examples

Overfitting in Decision Trees

Hypothesis space is complete! *Target function is surely in there*



No back tracking; *Greedy thus local minima*

Statistics-based search choices; *Robust to noisy data*

Inductive bias: heuristically prefer shortest tree

Decision trees will always overfit!

It is always possible to obtain zero training error on the input data with a deep enough tree (if there is no noise in the labels)



Possibly just noise, but the tree is grown larger to capture these examples


Avoiding Overfitting in Decision Trees

Pre-pruning/early stopping

- Typical stopping criterion
 - No error (if all instances belong to same class)
 - IF all the attribute values are same
- More restrictive conditions
 - Stop growing when data split is not statistically significant (example using chi-square test)
 - Stop if the number of instances is less than a predefined threshold
 - Stop if expanding does not significantly improve the measures (information gain)

Post-pruning after growing a full tree

- Separate data into training and validation sets
- Evaluate impact on validation set **when a node is “pruned”**
- **Greedily remove** node that improves performance the most
- Produces smallest version of most accurate subtree
- Typically use minimum description length (MDL) for post-pruning



Some Post-pruning Methods

Reduced-Error Pruning

- Use a validation set (tuning) to identify errors at every node
- Prune node with highest reduction error
- Repeat until error no longer reduces

Pessimistic Pruning

- No necessity of a validation set
- The error estimate at every node is conservative based on the training examples

Rule-post Pruning

- Convert tree to equivalent set of rules (how)?
- Prune each rule independently of others
- Sort final rules into desired sequence

Decision Trees

- **Decision Trees** – popular and a very efficient hypothesis space
 - Variable size: Any Boolean function can be represented
 - Handles discrete and continuous features
 - Handles classification **and regression**
 - Easy to implement
 - Easy to use
 - Computationally cheap
- Constructive **heuristic** search: built top-down by adding nodes
- **Decision trees will overfit!**
 - zero bias classifier (no mistakes) = large variance
 - must use tricks to find simpler trees
 - early stopping, pruning etc.